



Decision Processes

CISC 7404 - Decision Making

Steven Morad

University of Macau

Review	2
Markov Processes	3
Exercise	17
Markov Control Processes	19
Markov Decision Processes	26
Exercise	41
Coding	43
Exam Next Class	48

Review

Markov Processes

Markov Processes

Decisions must make some change in the world

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Before we look at decision making, let us think about how we model change in the world

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Before we look at decision making, let us think about how we model change in the world

Markov processes are a popular way to model the world

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Before we look at decision making, let us think about how we model change in the world

Markov processes are a popular way to model the world

Some things we can model using Markov processes:

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Before we look at decision making, let us think about how we model change in the world

Markov processes are a popular way to model the world

Some things we can model using Markov processes:

- Music

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Before we look at decision making, let us think about how we model change in the world

Markov processes are a popular way to model the world

Some things we can model using Markov processes:

- Music
- DNA sequences

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Before we look at decision making, let us think about how we model change in the world

Markov processes are a popular way to model the world

Some things we can model using Markov processes:

- Music
- DNA sequences
- Cryptography

Markov Processes

Decisions must make some change in the world

If they make no change, they do not matter, and are not decisions!

Before we look at decision making, let us think about how we model change in the world

Markov processes are a popular way to model the world

Some things we can model using Markov processes:

- Music
- DNA sequences
- Cryptography
- History

Markov Processes

We can model almost anything as a Markov process

Markov Processes

We can model almost anything as a Markov process

So what is a Markov process?

Markov Processes

We can model almost anything as a Markov process

So what is a Markov process?

It is a probabilistic model of dynamical systems that allows us to predict the future

Markov Processes

We can model almost anything as a Markov process

So what is a Markov process?

It is a probabilistic model of dynamical systems that allows us to predict the future

A Markov process consists of two parts

Markov Processes

We can model almost anything as a Markov process

So what is a Markov process?

It is a probabilistic model of dynamical systems that allows us to predict the future

A Markov process consists of two parts

The **state space**

S

Markov Processes

We can model almost anything as a Markov process

So what is a Markov process?

It is a probabilistic model of dynamical systems that allows us to predict the future

A Markov process consists of two parts

The **state space**

S

Outcome space describing the
state of our system

Markov Processes

We can model almost anything as a Markov process

So what is a Markov process?

It is a probabilistic model of dynamical systems that allows us to predict the future

A Markov process consists of two parts

The **state space**

$$S$$

Outcome space describing the
state of our system

The **state transition function**

$$\text{Tr} : S \mapsto \Delta S$$

Markov Processes

We can model almost anything as a Markov process

So what is a Markov process?

It is a probabilistic model of dynamical systems that allows us to predict the future

A Markov process consists of two parts

The **state space**

$$S$$

Outcome space describing the state of our system

The **state transition function**

$$\text{Tr} : S \mapsto \Delta S$$

$$\text{Tr}(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t)$$

$$s_{t+1} \sim \text{Tr}(\cdot \mid s_t)$$

Markov Processes

Problem: Predict the weather

Markov Processes

Problem: Predict the weather

$$S = \{\text{rain, cloud, sun}\} = \{R, C, S\}$$

Markov Processes

Problem: Predict the weather

$$S = \{\text{rain, cloud, sun}\} = \{R, C, S\}$$

$$\text{Tr}(s_{t+1} \mid s_t) = \text{Pr}(s_{t+1} \mid s_t)$$

Markov Processes

Problem: Predict the weather

$$S = \{\text{rain, cloud, sun}\} = \{R, C, S\}$$

$$\begin{aligned} \text{Tr}(s_{t+1} \mid s_t) &= \Pr(s_{t+1} \mid s_t) \\ &= \begin{bmatrix} \Pr(C \mid C) & \Pr(R \mid C) & \Pr(S \mid C) \\ \Pr(C \mid R) & \Pr(R \mid R) & \Pr(S \mid R) \\ \Pr(C \mid S) & \Pr(R \mid S) & \Pr(S \mid S) \end{bmatrix} \end{aligned}$$

Markov Processes

Problem: Predict the weather

$$S = \{\text{rain, cloud, sun}\} = \{R, C, S\}$$

$$\begin{aligned} \text{Tr}(s_{t+1} \mid s_t) &= \Pr(s_{t+1} \mid s_t) \\ &= \begin{bmatrix} \Pr(C \mid C) & \Pr(R \mid C) & \Pr(S \mid C) \\ \Pr(C \mid R) & \Pr(R \mid R) & \Pr(S \mid R) \\ \Pr(C \mid S) & \Pr(R \mid S) & \Pr(S \mid S) \end{bmatrix} \\ &= \begin{bmatrix} 0.4 & 0.3 & 0.3 \\ 0.5 & 0.3 & 0.2 \\ 0.5 & 0.1 & 0.4 \end{bmatrix} \end{aligned}$$

Markov Processes


Problem: Predict the weather

$$S = \{\text{rain, cloud, sun}\} = \{R, C, S\}$$

$$\text{Tr}(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t)$$

$$= \begin{bmatrix} \Pr(C \mid C) & \Pr(R \mid C) & \Pr(S \mid C) \\ \Pr(C \mid R) & \Pr(R \mid R) & \Pr(S \mid R) \\ \Pr(C \mid S) & \Pr(R \mid S) & \Pr(S \mid S) \end{bmatrix}$$

$$= \begin{bmatrix} 0.4 & 0.3 & 0.3 \\ 0.5 & 0.3 & 0.2 \\ 0.5 & 0.1 & 0.4 \end{bmatrix}$$

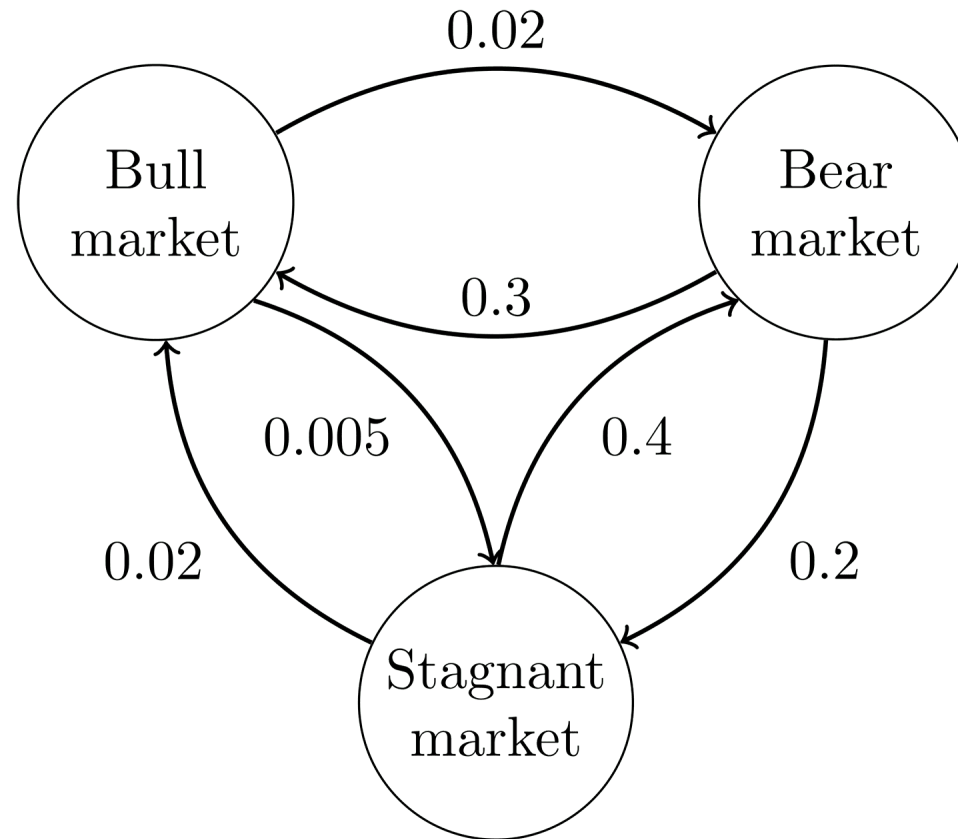


Markov Processes

We can model many other systems as Markov processes


Markov Processes

We can model many other systems as Markov processes



Markov Processes

Markov state diagram of a child behaviour



Markov Processes

Why is it called a **Markov** process?

Markov Processes

Why is it called a **Markov** process?

Markov property: The next state only depends on the current state

Markov Processes

Why is it called a **Markov** process?

Markov property: The next state only depends on the current state

$$\Pr(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_0)$$

Markov Processes

Why is it called a **Markov** process?

Markov property: The next state only depends on the current state

$$\Pr(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_0)$$

This is a very important condition we must always satisfy

Markov Processes

Why is it called a **Markov** process?

Markov property: The next state only depends on the current state

$$\Pr(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_0)$$

This is a very important condition we must always satisfy

If we cannot satisfy it, then the process is **not** Markov

Markov Processes

Why is it called a **Markov** process?

Markov property: The next state only depends on the current state

$$\Pr(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_0)$$

This is a very important condition we must always satisfy

If we cannot satisfy it, then the process is **not** Markov

$$\Pr(s_2 = \text{sun} \mid s_1 = \text{rain}, s_0 = \text{sun}) = 0.4$$

Markov Processes

Why is it called a **Markov** process?

Markov property: The next state only depends on the current state

$$\Pr(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_0)$$

This is a very important condition we must always satisfy

If we cannot satisfy it, then the process is **not** Markov

$$\Pr(s_2 = \text{sun} \mid s_1 = \text{rain}, s_0 = \text{sun}) = 0.4$$

$$\Pr(s_2 = \text{sun} \mid s_1 = \text{rain}) = 0.3$$

Markov Processes

Why is it called a **Markov** process?

Markov property: The next state only depends on the current state

$$\Pr(s_{t+1} \mid s_t) = \Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_0)$$

This is a very important condition we must always satisfy

If we cannot satisfy it, then the process is **not** Markov

$$\Pr(s_2 = \text{sun} \mid s_1 = \text{rain}, s_0 = \text{sun}) = 0.4$$

$$\Pr(s_2 = \text{sun} \mid s_1 = \text{rain}) = 0.3$$

$0.3 \neq 0.4$, $\Pr(s_{t+1} \mid s_t) \neq \Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_0)$, **not** Markov

Markov Processes

We can visualize the Markov property too

Markov Processes


We can visualize the Markov property too

To compute the next node, we only look at the current node

Markov Processes

We can visualize the Markov property too

To compute the next node, we only look at the current node



Markov Processes

We can predict the future using Markov processes

Markov Processes

We can predict the future using Markov processes

Chain transition probabilities together to estimate s_t

Markov Processes

We can predict the future using Markov processes

Chain transition probabilities together to estimate s_t

$$\Pr(s_1 \mid s_0 = s)$$

Markov Processes

We can predict the future using Markov processes

Chain transition probabilities together to estimate s_t

$$\Pr(s_1 \mid s_0 = s)$$

$$\Pr(s_2 \mid s_0 = s) = \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Markov Processes

We can predict the future using Markov processes

Chain transition probabilities together to estimate s_t

$$\Pr(s_1 \mid s_0 = s)$$

$$\Pr(s_2 \mid s_0 = s) = \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$



Paths from all possible s_1 to s_2

Markov Processes

We can predict the future using Markov processes

Chain transition probabilities together to estimate s_t

$$\Pr(s_1 \mid s_0 = s)$$

$$\Pr(s_2 \mid s_0 = s) = \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$



Paths from all possible s_1 to s_2

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \Pr(s_3 \mid s_2) \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Markov Processes

We can predict the future using Markov processes

Chain transition probabilities together to estimate s_t

$$\Pr(s_1 \mid s_0 = s)$$

$$\Pr(s_2 \mid s_0 = s) = \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$



Paths from all possible s_1 to s_2

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \Pr(s_3 \mid s_2) \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Can we derive a general form for $P(s_n \mid s_0)$?

Markov Processes

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \Pr(s_3 \mid s_2) \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Markov Processes

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \Pr(s_3 \mid s_2) \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Product of sum is the sum of products

Markov Processes

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \Pr(s_3 \mid s_2) \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Product of sum is the sum of products, move the sum outside

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \sum_{s_1 \in S} \Pr(s_3 \mid s_2) \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Markov Processes

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \Pr(s_3 \mid s_2) \sum_{s_1 \in S} \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Product of sum is the sum of products, move the sum outside

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_2 \in S} \sum_{s_1 \in S} \Pr(s_3 \mid s_2) \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Combine the sums

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_1, s_2 \in S} \Pr(s_3 \mid s_2) \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Markov Processes

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_1, s_2 \in S} \Pr(s_3 \mid s_2) \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Markov Processes

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_1, s_2 \in S} \Pr(s_3 \mid s_2) \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Combine the products

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_1, s_2 \in S} \prod_{t=0}^2 \Pr(s_{t+1} \mid s_t)$$

Markov Processes

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_1, s_2 \in S} \Pr(s_3 \mid s_2) \Pr(s_2 \mid s_1) \Pr(s_1 \mid s_0 = s)$$

Combine the products

$$\Pr(s_3 \mid s_0 = s) = \sum_{s_1, s_2 \in S} \prod_{t=0}^2 \Pr(s_{t+1} \mid s_t)$$

Generalize to any timestep n

$$\Pr(s_n \mid s_0) = \sum_{s_1, s_2, \dots, s_{n-1} \in S} \prod_{t=0}^{n-1} \Pr(s_{t+1} \mid s_t)$$

Markov Processes

$$\Pr(s_n \mid s_0) = \sum_{s_1, s_2, \dots, s_{n-1} \in S} \prod_{t=0}^{n-1} \Pr(s_{t+1} \mid s_t)$$

Markov Processes

$$\Pr(s_n \mid s_0) = \sum_{s_1, s_2, \dots, s_{n-1} \in S} \prod_{t=0}^{n-1} \Pr(s_{t+1} \mid s_t)$$

This expression tells us how the Markov process evolves over time

Markov Processes

$$\Pr(s_n \mid s_0) = \sum_{s_1, s_2, \dots, s_{n-1} \in S} \prod_{t=0}^{n-1} \Pr(s_{t+1} \mid s_t)$$

This expression tells us how the Markov process evolves over time

We can predict the future, n timesteps from now

Markov Processes

$$\Pr(s_n \mid s_0) = \sum_{s_1, s_2, \dots, s_{n-1} \in S} \prod_{t=0}^{n-1} \Pr(s_{t+1} \mid s_t)$$

This expression tells us how the Markov process evolves over time

We can predict the future, n timesteps from now

If s is the state of the world, you can predict the future of the world

Markov Processes

$$\Pr(s_n \mid s_0) = \sum_{s_1, s_2, \dots, s_{n-1} \in S} \prod_{t=0}^{n-1} \Pr(s_{t+1} \mid s_t)$$

This expression tells us how the Markov process evolves over time

We can predict the future, n timesteps from now

If s is the state of the world, you can predict the future of the world

If s represents someone's mind, you can predict their future thoughts

Markov Processes

We can predict a state s_n , but a Markov process never ends

Markov Processes

We can predict a state s_n , but a Markov process never ends

The future infinite, there is always a next state

Markov Processes

We can predict a state s_n , but a Markov process never ends

The future infinite, there is always a next state

However, many processes we like to model eventually end

Markov Processes

We can predict a state s_n , but a Markov process never ends

The future infinite, there is always a next state

However, many processes we like to model eventually end

- Dying in a video game

Markov Processes

We can predict a state s_n , but a Markov process never ends

The future infinite, there is always a next state

However, many processes we like to model eventually end

- Dying in a video game
- Completing a task

Markov Processes

We can predict a state s_n , but a Markov process never ends

The future infinite, there is always a next state

However, many processes we like to model eventually end

- Dying in a video game
- Completing a task
- Running out of money

Markov Processes

We can predict a state s_n , but a Markov process never ends

The future infinite, there is always a next state

However, many processes we like to model eventually end

- Dying in a video game
- Completing a task
- Running out of money

Question: How can we model a Markov process that ends?

Markov Processes

We can predict a state s_n , but a Markov process never ends

The future infinite, there is always a next state

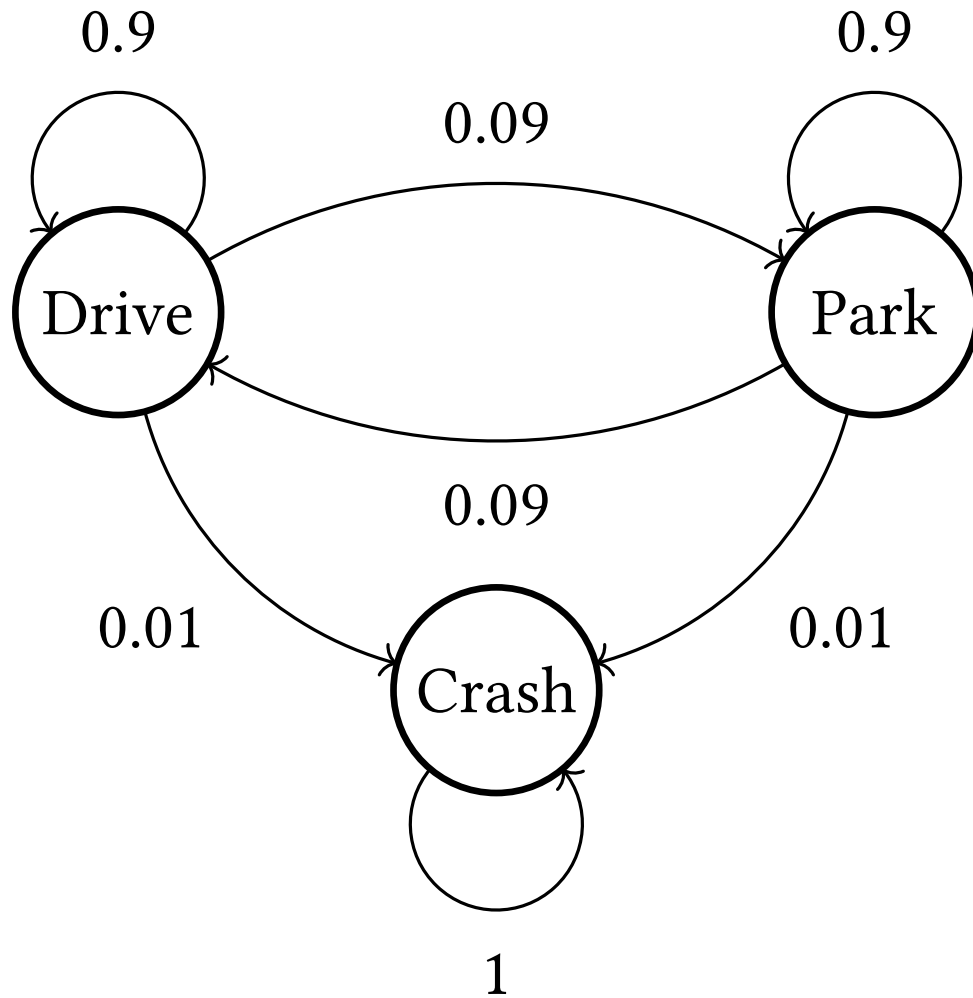
However, many processes we like to model eventually end

- Dying in a video game
- Completing a task
- Running out of money

Question: How can we model a Markov process that ends?


Answer: We create a **terminal state** that we can enter but cannot leave

Markov Processes




Markov Processes

Upon reaching a terminal state, we get stuck




Markov Processes



Upon reaching a terminal state, we get stuck

Once we crash our car, we cannot drive or park any more

Markov Processes



Upon reaching a terminal state, we get stuck

Once we crash our car, we cannot drive or park any more

The only transition from a terminal state is back to itself

$$\Pr(s_{t+1} = \text{term} \mid s_t = \text{term}) = 1$$

$$\Pr(s_{t+1} = \text{not term} \mid s_t = \text{term}) = 0$$

Exercise

Exercise

Design an Markov process about a problem you care about

Exercise

Design an Markov process about a problem you care about

- 4 states

Exercise

Design an Markov process about a problem you care about

- 4 states
- State transition function $\text{Tr} = \Pr(s_{t+1} \mid s_t)$ for all $s_t, s_{t+1} \in S$

Exercise

Design an Markov process about a problem you care about

- 4 states
- State transition function $\text{Tr} = \Pr(s_{t+1} \mid s_t)$ for all $s_t, s_{t+1} \in S$
- Create a terminal state

Exercise

Design an Markov process about a problem you care about

- 4 states
- State transition function $\text{Tr} = \Pr(s_{t+1} \mid s_t)$ for all $s_t, s_{t+1} \in S$
- Create a terminal state
- Given a starting state s_0 , what will your state distribution be for s_2 ?

Exercise

Design an Markov process about a problem you care about

- 4 states
- State transition function $\text{Tr} = \Pr(s_{t+1} \mid s_t)$ for all $s_t, s_{t+1} \in S$
- Create a terminal state
- Given a starting state s_0 , what will your state distribution be for s_2 ?

$$\Pr(s_n \mid s_0) = \sum_{s_1, s_2, \dots, s_{n-1} \in S} \prod_{t=0}^{n-1} \Pr(s_{t+1} \mid s_t)$$

Markov Control Processes

Markov Control Processes

Markov processes model complex evolving processes

Markov Control Processes

Markov processes model complex evolving processes

But this course is on decision making, how can we model decision making in a Markov process?

We can't

Markov Control Processes

Markov processes model complex evolving processes

But this course is on decision making, how can we model decision making in a Markov process?

We can't

Markov processes follow the state transition function T

Markov Control Processes

Markov processes model complex evolving processes

But this course is on decision making, how can we model decision making in a Markov process?

We can't

Markov processes follow the state transition function T

The future of the system is already determined

Markov Control Processes

Markov processes model complex evolving processes

But this course is on decision making, how can we model decision making in a Markov process?

We can't

Markov processes follow the state transition function T

The future of the system is already determined

There is no room for decisions to change the fate of the system

Markov Control Processes

Markov processes model complex evolving processes

But this course is on decision making, how can we model decision making in a Markov process?

We can't

Markov processes follow the state transition function T

The future of the system is already determined

There is no room for decisions to change the fate of the system

We will modify the Markov process for decision making

Markov Control Processes

Markov processes model complex evolving processes

But this course is on decision making, how can we model decision making in a Markov process?

We can't

Markov processes follow the state transition function T

The future of the system is already determined

There is no room for decisions to change the fate of the system

We will modify the Markov process for decision making

The point of decision making is to choose our fate

Markov Control Processes

A Markov process models the predetermined evolution of some system

Markov Control Processes

A Markov process models the predetermined evolution of some system

We call this system the **environment**, because we cannot control it

Markov Control Processes

A Markov process models the predetermined evolution of some system

We call this system the **environment**, because we cannot control it

The **agent** lives in the environment

Markov Control Processes

A Markov process models the predetermined evolution of some system

We call this system the **environment**, because we cannot control it

The **agent** lives in the environment

The agent makes decisions

Markov Control Processes

A Markov process models the predetermined evolution of some system

We call this system the **environment**, because we cannot control it

The **agent** lives in the environment

The agent makes decisions

The agent changes the environment with its decisions

Markov Control Processes

The agent takes **actions** $a \in A$ that change the environment

Markov Control Processes

The agent takes **actions** $a \in A$ that change the environment

The action space A defines what our agent can do

Markov Control Processes

The agent takes **actions** $a \in A$ that change the environment

The action space A defines what our agent can do

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov Control Processes

The agent takes **actions** $a \in A$ that change the environment

The action space A defines what our agent can do

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov control process

$$(S, A, \text{Tr})$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

Markov Control Processes

The agent takes **actions** $a \in A$ that change the environment

The action space A defines what our agent can do

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov control process

$$(S, A, \text{Tr})$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

In a Markov process, the future follows a predefined evolution

Markov Control Processes

The agent takes **actions** $a \in A$ that change the environment

The action space A defines what our agent can do

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov control process

$$(S, A, \text{Tr})$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

In a Markov process, the future follows a predefined evolution

In a Markov control process, we can control the evolution!

Markov Control Processes

The agent takes **actions** $a \in A$ that change the environment

The action space A defines what our agent can do

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov control process

$$(S, A, \text{Tr})$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

In a Markov process, the future follows a predefined evolution

In a Markov control process, we can control the evolution!

Let us see an example

Markov Control Processes

$$S = \{\text{Healthy}, \text{Sick}, \text{Dead}\}$$

Markov Control Processes


$$S = \{\text{Healthy}, \text{Sick}, \text{Dead}\}$$

$$A = \{\text{Nothing}, \text{Medicine}\} = \{N, M\}$$


Markov Control Processes

$S = \{\text{Healthy}, \text{Sick}, \text{Dead}\}$

$A = \{\text{Nothing}, \text{Medicine}\} = \{N, M\}$




Markov Control Processes




Markov Control Processes

The **trajectory** contains the states and actions until a terminal state




Markov Control Processes



The **trajectory** contains the states and actions until a terminal state

$$\tau = \begin{bmatrix} s_0 & a_0 \\ s_1 & a_1 \\ s_2 & a_2 \\ \vdots & \vdots \\ s_{n-1} & a_{n-1} \\ s_n & \emptyset \end{bmatrix}$$


Markov Control Processes



The **trajectory** contains the states and actions until a terminal state

$$\tau = \begin{bmatrix} s_0 & a_0 \\ s_1 & a_1 \\ s_2 & a_2 \\ \vdots & \vdots \\ s_{n-1} & a_{n-1} \\ s_n & \emptyset \end{bmatrix} = \begin{bmatrix} \text{Healthy} & \text{Nothing} \\ \text{Sick} & \text{Nothing} \\ \text{Sick} & \text{Medicine} \\ \vdots & \vdots \\ \text{Sick} & \text{Nothing} \\ \text{Dead} & \emptyset \end{bmatrix}$$

Markov Control Processes



The **trajectory** contains the states and actions until a terminal state

$$\tau = \begin{bmatrix} s_0 & a_0 \\ s_1 & a_1 \\ s_2 & a_2 \\ \vdots & \vdots \\ s_{n-1} & a_{n-1} \\ s_n & \emptyset \end{bmatrix} = \begin{bmatrix} \text{Healthy} & \text{Nothing} \\ \text{Sick} & \text{Nothing} \\ \text{Sick} & \text{Medicine} \\ \vdots & \vdots \\ \text{Sick} & \text{Nothing} \\ \text{Dead} & \emptyset \end{bmatrix}$$

If there is no terminal state, the trajectory can be infinitely long!

Markov Control Processes

Markov control processes let us control which states we visit

Markov Control Processes

Markov control processes let us control which states we visit

They do not tell us which states are good to visit, or provide a goal

Markov Control Processes

Markov control processes let us control which states we visit

They do not tell us which states are good to visit, or provide a goal

How can we make optimal decisions if we do not have a goal or objective?

Markov Control Processes

Markov control processes let us control which states we visit

They do not tell us which states are good to visit, or provide a goal

How can we make optimal decisions if we do not have a goal or objective?

We need a way to determine “good” and “bad” decisions

Markov Decision Processes

Markov Decision Processes

Markov decision processes (MDPs) add a measure of “goodness” to Markov control processes

Markov Decision Processes

Markov decision processes (MDPs) add a measure of “goodness” to Markov control processes

We use a **reward function** R to measure the goodness of a specific state

Markov Decision Processes

Markov decision processes (MDPs) add a measure of “goodness” to Markov control processes

We use a **reward function** R to measure the goodness of a specific state

Sutton and Barto:

$$R : S \times A \mapsto \mathbb{R}$$

Markov Decision Processes

Markov decision processes (MDPs) add a measure of “goodness” to Markov control processes

We use a **reward function** R to measure the goodness of a specific state

Sutton and Barto:

$$R : S \times A \mapsto \mathbb{R}$$

Other books:

$$R : S \times A \times S \mapsto \mathbb{R}$$

Markov Decision Processes

Markov decision processes (MDPs) add a measure of “goodness” to Markov control processes

We use a **reward function** R to measure the goodness of a specific state

Sutton and Barto:

$$R : S \times A \mapsto \mathbb{R}$$

Other books:

$$R : S \times A \times S \mapsto \mathbb{R}$$

This course:

$$R : S \mapsto \mathbb{R}$$

Markov Decision Processes

Markov decision processes (MDPs) add a measure of “goodness” to Markov control processes

We use a **reward function** R to measure the goodness of a specific state

Sutton and Barto:

$$R : S \times A \mapsto \mathbb{R}$$

Other books:

$$R : S \times A \times S \mapsto \mathbb{R}$$

This course:

$$R : S \mapsto \mathbb{R}$$

For now, I will use the simplest one

Markov Decision Processes

Markov decision processes (MDPs) add a measure of “goodness” to Markov control processes

We use a **reward function** R to measure the goodness of a specific state

Sutton and Barto:

$$R : S \times A \mapsto \mathbb{R}$$

Other books:

$$R : S \times A \times S \mapsto \mathbb{R}$$

This course:

$$R : S \mapsto \mathbb{R}$$

For now, I will use the simplest one

You can always make these equivalent by modifying the MDP

Markov Decision Processes

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov Decision Processes

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov control
process

$$(S, A, \text{Tr})$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

Markov Decision Processes

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov control
process

$$(S, A, \text{Tr})$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

Markov decision
process

$$(S, A, \text{Tr}, R, \gamma)$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

$$R : S \mapsto \mathbb{R}$$

In an MDP, an **episode** contains
the trajectory and also the rewards

Markov Decision Processes

Markov process

$$(S, \text{Tr})$$

$$\text{Tr} : S \mapsto \Delta S$$

Markov control
process

$$(S, A, \text{Tr})$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

Markov decision
process

$$(S, A, \text{Tr}, R, \gamma)$$

$$\text{Tr} : S \times A \mapsto \Delta S$$

$$R : S \mapsto \mathbb{R}$$

In an MDP, an **episode** contains
the trajectory and also the rewards

$$E = \begin{bmatrix} s_0 & a_0 & r_0 \\ s_1 & a_1 & r_1 \\ \vdots & \vdots & \vdots \\ s_{n-1} & a_{n-1} & r_{n-1} \\ s_n & \emptyset & \emptyset \end{bmatrix} = [\boldsymbol{\tau} \quad \boldsymbol{r}]$$

Markov Decision Processes

Reward is good, we want to maximize the reward

Markov Decision Processes

Reward is good, we want to maximize the reward

The reward function determines the agent behavior

Markov Decision Processes

Reward is good, we want to maximize the reward

The reward function determines the agent behavior

$s_d = \text{Dumpling}$

$s_n = \text{Noodle}$

Markov Decision Processes

Reward is good, we want to maximize the reward

The reward function determines the agent behavior

$s_d = \text{Dumpling}$

$s_n = \text{Noodle}$

$$R(s_d) = 10$$

$$R(s_n) = 15$$

Markov Decision Processes

Reward is good, we want to maximize the reward

The reward function determines the agent behavior

$s_d = \text{Dumpling}$

$s_n = \text{Noodle}$

$$R(s_d) = 10$$

$$R(s_n) = 15$$

Result: Eat noodle

Markov Decision Processes

Reward is good, we want to maximize the reward

The reward function determines the agent behavior

$s_d = \text{Dumpling}$

$s_n = \text{Noodle}$

$$R(s_d) = 10$$

$$R(s_n) = 15$$

Result: Eat noodle

$$R(s_d) = 5$$

$$R(s_n) = -3$$

Markov Decision Processes

Reward is good, we want to maximize the reward

The reward function determines the agent behavior

s_d = Dumpling

s_n = Noodle

$$R(s_d) = 10$$

$$R(s_n) = 15$$

Result: Eat noodle

$$R(s_d) = 5$$

$$R(s_n) = -3$$

Result: Eat dumpling

Markov Decision Processes

Reward is good, we want to maximize the reward

The reward function determines the agent behavior

$s_d = \text{Dumpling}$

$s_n = \text{Noodle}$

$$R(s_d) = 10$$

$$R(s_n) = 15$$

Result: Eat noodle

$$R(s_d) = 5$$

$$R(s_n) = -3$$

Result: Eat dumpling

We can write this mathematically as

$$\arg \max_{s \in S} R(s)$$

Markov Decision Processes

However, maximizing the reward is not always ideal

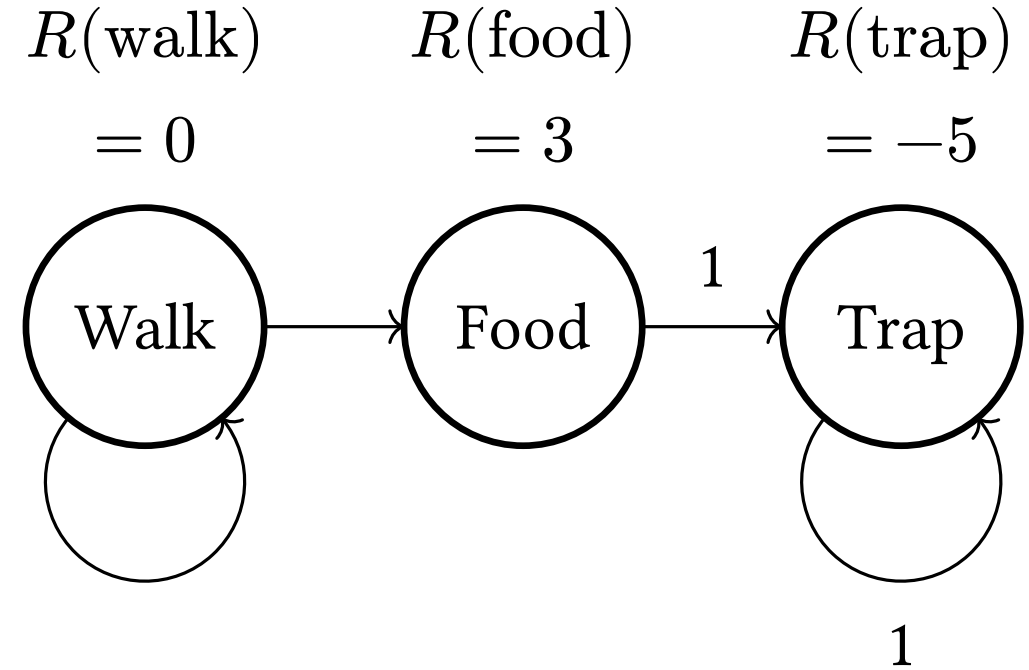
Markov Decision Processes

However, maximizing the reward is not always ideal



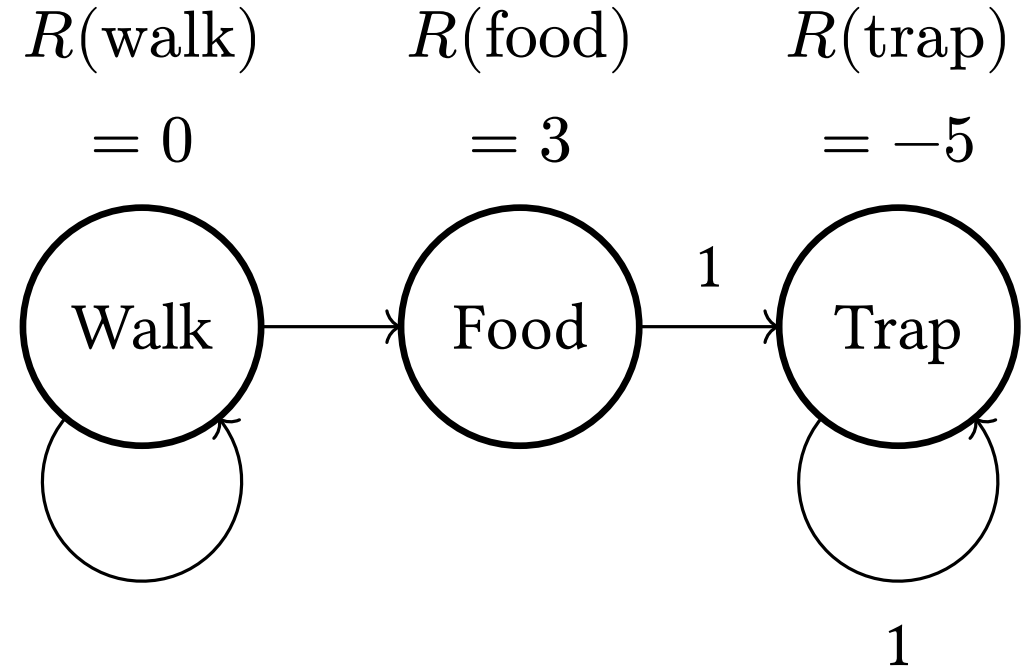

Markov Decision Processes

However, maximizing the reward is not always ideal



Markov Decision Processes

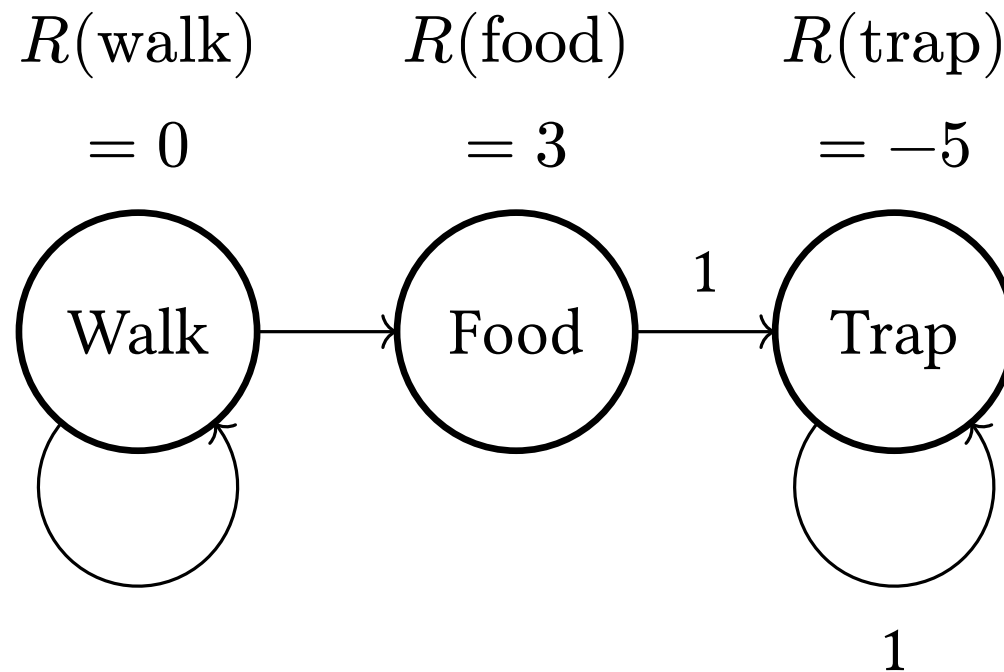

However, maximizing the reward is not always ideal



$$\arg \max_{a \in A} R(s) = \text{take the food}$$

Markov Decision Processes

However, maximizing the reward is not always ideal



$$\arg \max_{a \in A} R(s) = \text{take the food}$$

If we maximize the reward, we are **too greedy**

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Instead, we maximize the **cumulative sum** of rewards

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Instead, we maximize the **cumulative sum** of rewards

We think about how our actions now will impact reward the future

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Instead, we maximize the **cumulative sum** of rewards

We think about how our actions now will impact reward the future

We call the cumulative sum of rewards, the **return**

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Instead, we maximize the **cumulative sum** of rewards

We think about how our actions now will impact reward the future

We call the cumulative sum of rewards, the **return**

$$G : \mathbb{R}^n \mapsto \mathbb{R}$$

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Instead, we maximize the **cumulative sum** of rewards

We think about how our actions now will impact reward the future

We call the cumulative sum of rewards, the **return**

$$G : \mathbb{R}^n \mapsto \mathbb{R}$$

$$G : S^n \times A^{n-1} \mapsto \mathbb{R}$$

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Instead, we maximize the **cumulative sum** of rewards

We think about how our actions now will impact reward the future

We call the cumulative sum of rewards, the **return**

$$G : \mathbb{R}^n \mapsto \mathbb{R}$$

$$G : S^n \times A^{n-1} \mapsto \mathbb{R}$$

$$G(r_0, r_1, \dots) = \sum_{t=0}^{\infty} r_t$$

Markov Decision Processes

Maximizing the immediate reward can result in bad agents

Instead, we maximize the **cumulative sum** of rewards

We think about how our actions now will impact reward the future

We call the cumulative sum of rewards, the **return**

$$G : \mathbb{R}^n \mapsto \mathbb{R}$$

$$G(r_0, r_1, \dots) = \sum_{t=0}^{\infty} r_t$$

$$G : S^n \times A^{n-1} \mapsto \mathbb{R}$$

$$\begin{aligned} G(\tau) &= G(s_0, a_0, s_1, a_1, \dots) \\ &= \sum_{t=0}^{\infty} R(s_{t+1}) \end{aligned}$$


Markov Decision Processes

$R(\text{walk})$ $R(\text{food})$ $R(\text{trap})$

$= 0$

$= 3$

$= -5$



$$G(\tau_{\text{greedy}}) = R(\text{food}) + R(\text{trap}) + R(\text{trap}) + \dots = 3 - 5 - 5 - \dots = -\infty$$


Markov Decision Processes

$R(\text{walk})$ $R(\text{food})$ $R(\text{trap})$

$= 0$

$= 3$

$= -5$



$$G(\tau_{\text{greedy}}) = R(\text{food}) + R(\text{trap}) + R(\text{trap}) + \dots = 3 - 5 - 5 - \dots = -\infty$$

$$G(\tau_{\text{smart}}) = R(\text{walk}) + R(\text{walk}) + R(\text{walk}) + \dots = 0 + 0 + \dots = 0$$


Markov Decision Processes

$R(\text{walk})$ $R(\text{food})$ $R(\text{trap})$

$= 0$

$= 3$

$= -5$



$$G(\tau_{\text{greedy}}) = R(\text{food}) + R(\text{trap}) + R(\text{trap}) + \dots = 3 - 5 - 5 - \dots = -\infty$$

$$G(\tau_{\text{smart}}) = R(\text{walk}) + R(\text{walk}) + R(\text{walk}) + \dots = 0 + 0 + \dots = 0$$

By considering the future rewards, we can make optimal decisions

Markov Decision Processes

Consider one more example

Markov Decision Processes


Consider one more example

$R(\text{walk})$ $R(\text{food})$ $R(\text{sleep})$

$= 0$

$= 3$


$= 0$



Markov Decision Processes

Consider one more example

$R(\text{walk}) = 0$ $R(\text{food}) = 3$ $R(\text{sleep}) = 0$




Question: What is the optimal sequence of states?

Markov Decision Processes

Consider one more example

$$\begin{array}{ccc} R(\text{walk}) & R(\text{food}) & R(\text{sleep}) \\ = 0 & = 3 & = 0 \end{array}$$




Question: What is the optimal sequence of states?

$$\text{Walk} + \text{Food} + \text{Sleep} + \dots = 0 + 3 + 0 + \dots = 3$$

Markov Decision Processes

Consider one more example

$$\begin{array}{ccc} R(\text{walk}) & R(\text{food}) & R(\text{sleep}) \\ = 0 & = 3 & = 0 \end{array}$$



Question: What is the optimal sequence of states?

$$\text{Walk} + \text{Food} + \text{Sleep} + \dots = 0 + 3 + 0 + \dots = 3$$

$$\text{Walk} + \text{Walk} + \dots + \text{Food} + \text{Sleep} + \dots = 0 + 0 + \dots + 3 + 0 + \dots = 3$$

Markov Decision Processes

The return is a sum

$$G(\boldsymbol{\tau}) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Markov Decision Processes

The return is a sum

$$G(\boldsymbol{\tau}) = \sum_{t=0}^{\infty} R(s_{t+1})$$

We can eat food now, or in 1000 years, the return is the same

Markov Decision Processes

The return is a sum

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

We can eat food now, or in 1000 years, the return is the same

Question: Is this how humans make decisions?

Markov Decision Processes

The return is a sum

$$G(\boldsymbol{\tau}) = \sum_{t=0}^{\infty} R(s_{t+1})$$

We can eat food now, or in 1000 years, the return is the same

Question: Is this how humans make decisions?

Answer: No, humans are a little bit greedy

Markov Decision Processes

The return is a sum

$$G(\boldsymbol{\tau}) = \sum_{t=0}^{\infty} R(s_{t+1})$$

We can eat food now, or in 1000 years, the return is the same

Question: Is this how humans make decisions?

Answer: No, humans are a little bit greedy

Experiment: Place a cookie in front of a child. If they do not eat the cookie for 5 minutes, they get two cookies

Markov Decision Processes

The return is a sum

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

We can eat food now, or in 1000 years, the return is the same

Question: Is this how humans make decisions?

Answer: No, humans are a little bit greedy

Experiment: Place a cookie in front of a child. If they do not eat the cookie for 5 minutes, they get two cookies

Question: What happens?

Markov Decision Processes

The return is a sum

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

We can eat food now, or in 1000 years, the return is the same

Question: Is this how humans make decisions?

Answer: No, humans are a little bit greedy

Experiment: Place a cookie in front of a child. If they do not eat the cookie for 5 minutes, they get two cookies

Question: What happens?

Answer: Child eats the cookie immediately

Markov Decision Processes

The return is a sum

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

We can eat food now, or in 1000 years, the return is the same

Question: Is this how humans make decisions?

Answer: No, humans are a little bit greedy

Experiment: Place a cookie in front of a child. If they do not eat the cookie for 5 minutes, they get two cookies

Question: What happens?

Answer: Child eats the cookie immediately

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Question: How can we modify the return to prefer rewards sooner?

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Question: How can we modify the return to prefer rewards sooner?

What if we make future rewards less important?

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Question: How can we modify the return to prefer rewards sooner?

What if we make future rewards less important?

$$R(s_{t+1}) = \{1 \mid s_{t+1} \in S\}$$

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Question: How can we modify the return to prefer rewards sooner?

What if we make future rewards less important?

$$R(s_{t+1}) = \{1 \mid s_{t+1} \in S\} \qquad G(\tau) = \sum_{t=0}^{\infty} 1 = 1 + 1 + \dots$$

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Question: How can we modify the return to prefer rewards sooner?

What if we make future rewards less important?

$$R(s_{t+1}) = \{1 \mid s_{t+1} \in S\}$$

$$G(\tau) = \sum_{t=0}^{\infty} 1 = 1 + 1 + \dots$$

$$G(\tau) = \quad ? \quad = 1 + 0.9 + 0.8 + \dots$$

Markov Decision Processes

Timing matters, humans prefer reward sooner rather than later

$$G(\tau) = \sum_{t=0}^{\infty} R(s_{t+1})$$

Question: How can we modify the return to prefer rewards sooner?

What if we make future rewards less important?

$$R(s_{t+1}) = \{1 \mid s_{t+1} \in S\}$$

$$G(\tau) = \sum_{t=0}^{\infty} 1 = 1 + 1 + \dots$$

$$G(\tau) = \quad ? \quad = 1 + 0.9 + 0.8 + \dots$$

Question: How?

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

With $\gamma = 0$

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

$$G(\boldsymbol{\tau}) = \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

With $\gamma = 0$

$$G(\boldsymbol{\tau}) = 1 + 1 + 1 + \dots$$

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

With $\gamma = 0$

$$G(\tau) = 1 + 1 + 1 + \dots$$

With $\gamma = 0.9$

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

$$G(\boldsymbol{\tau}) = \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

With $\gamma = 0$

$$G(\boldsymbol{\tau}) = 1 + 1 + 1 + \dots$$

With $\gamma = 0.9$

$$G(\boldsymbol{\tau}) = (0.9^0 \cdot 1) + (0.9^1 \cdot 1) + (0.9^2 \cdot 1) + \dots$$

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

With $\gamma = 0$

$$G(\tau) = 1 + 1 + 1 + \dots$$

With $\gamma = 0.9$

$$G(\tau) = (0.9^0 \cdot 1) + (0.9^1 \cdot 1) + (0.9^2 \cdot 1) + \dots$$

$$G(\tau) = 1 + 0.9 + 0.81 + \dots$$

We call this the **discounted return**

Markov Decision Processes

We can introduce a **discount** term $\gamma \in [0, 1]$ to the return

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

With $\gamma = 0$

$$G(\tau) = 1 + 1 + 1 + \dots$$

With $\gamma = 0.9$

$$G(\tau) = (0.9^0 \cdot 1) + (0.9^1 \cdot 1) + (0.9^2 \cdot 1) + \dots$$

$$G(\tau) = 1 + 0.9 + 0.81 + \dots$$

We call this the **discounted return**

The discounted return lets makes us prefer rewards sooner, like humans

Markov Decision Processes

For the rest of the course, we maximize the discounted return

$$\arg \max_{\tau} G(\tau) = \arg \max_{s \in S} \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

Markov Decision Processes

For the rest of the course, we maximize the discounted return

$$\arg \max_{\tau} G(\tau) = \arg \max_{s \in S} \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

If our agent maximizes the discounted return, then it is **optimal**

Markov Decision Processes

Let us review

Markov Decision Processes

Let us review

Definition: A Markov decision process (MDP) is a tuple (S, A, T, R, γ)

Markov Decision Processes

Let us review

Definition: A Markov decision process (MDP) is a tuple (S, A, T, R, γ)

- S is the state space

Markov Decision Processes

Let us review

Definition: A Markov decision process (MDP) is a tuple (S, A, T, R, γ)

- S is the state space
- A is the action space

Markov Decision Processes

Let us review

Definition: A Markov decision process (MDP) is a tuple (S, A, T, R, γ)

- S is the state space
- A is the action space
- $\text{Tr} : S \times A \mapsto \Delta S$ is the state transition function

Markov Decision Processes

Let us review

Definition: A Markov decision process (MDP) is a tuple (S, A, T, R, γ)

- S is the state space
- A is the action space
- $\text{Tr} : S \times A \mapsto \Delta S$ is the state transition function
- $R : S \mapsto \mathbb{R}$ is the reward function

Markov Decision Processes

Let us review

Definition: A Markov decision process (MDP) is a tuple (S, A, T, R, γ)

- S is the state space
- A is the action space
- $\text{Tr} : S \times A \mapsto \Delta S$ is the state transition function
- $R : S \mapsto \mathbb{R}$ is the reward function
- $\gamma \in [0, 1]$ is the discount factor

Markov Decision Processes

Let us review

Definition: A Markov decision process (MDP) is a tuple (S, A, T, R, γ)

- S is the state space
- A is the action space
- $\text{Tr} : S \times A \mapsto \Delta S$ is the state transition function
- $R : S \mapsto \mathbb{R}$ is the reward function
- $\gamma \in [0, 1]$ is the discount factor

Markov Decision Processes

Reinforcement learning is designed to solve MDPs

Markov Decision Processes

Reinforcement learning is designed to solve MDPs

In reinforcement learning, we have a single goal

Markov Decision Processes

Reinforcement learning is designed to solve MDPs

In reinforcement learning, we have a single goal

Maximize the discounted return of the MDP

Markov Decision Processes

Reinforcement learning is designed to solve MDPs

In reinforcement learning, we have a single goal

Maximize the discounted return of the MDP

$$\arg \max_{\tau} G(\tau) = \arg \max_{s \in S} \sum_{t=0}^{\infty} \gamma^t R(s_{t+1})$$

You must understand the discounted return!

Markov Decision Processes

Understanding MDPs is the **most important part** of RL

Markov Decision Processes

Understanding MDPs is the **most important part** of RL

Existing software can train RL agents on your MDP

Markov Decision Processes

Understanding MDPs is the **most important part** of RL

Existing software can train RL agents on your MDP

You can train an RL agent without understanding RL

Markov Decision Processes

Understanding MDPs is the **most important part** of RL

Existing software can train RL agents on your MDP

You can train an RL agent without understanding RL

You can only train an agent if you can model your problem as an MDP

Markov Decision Processes

Understanding MDPs is the **most important part** of RL

Existing software can train RL agents on your MDP

You can train an RL agent without understanding RL

You can only train an agent if you can model your problem as an MDP

Make sure you understand MDPs!

Exercise

Exercise

Design a Super Mario Bros MDP



Exercise

Design a Super Mario Bros MDP

- Reward function R



Exercise

Design a Super Mario Bros MDP

- Reward function R
- Discount factor γ



Exercise



Design a Super Mario Bros MDP

- Reward function R
- Discount factor γ

Your states are: eat mushroom, collect coins, die, game over

Exercise



Design a Super Mario Bros MDP

- Reward function R
- Discount factor γ

Your states are: eat mushroom, collect coins, die, game over

Compute discounted return for:

Exercise



Design a Super Mario Bros MDP

- Reward function R
- Discount factor γ

Your states are: eat mushroom, collect coins, die, game over

Compute discounted return for:

- Eat mushroom at $t = 10$

Exercise



Design a Super Mario Bros MDP

- Reward function R
- Discount factor γ

Your states are: eat mushroom, collect coins, die, game over

Compute discounted return for:

- Eat mushroom at $t = 10$
- Collect coins at $t = 11, 12$

Exercise



Design a Super Mario Bros MDP

- Reward function R
- Discount factor γ

Your states are: eat mushroom, collect coins, die, game over

Compute discounted return for:

- Eat mushroom at $t = 10$
- Collect coins at $t = 11, 12$
- Die to bowser at $t = 20$
- Game over screen at $t = 21 \dots \infty$
- $r = 0$ for other timesteps

Coding

Coding

In this course, we will implemented MDPs using **gymnasium**

Coding

In this course, we will implemented MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Coding

In this course, we will implement MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Gymnasium provides an **environment** (MDP) API

Coding

In this course, we will implemented MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Gymnasium provides an **environment** (MDP) API

Must define:

Coding

In this course, we will implemented MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Gymnasium provides an **environment** (MDP) API

Must define:

- state space (S)

Coding

In this course, we will implemented MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Gymnasium provides an **environment** (MDP) API

Must define:

- state space (S)
- action space (A)

Coding

In this course, we will implemented MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Gymnasium provides an **environment** (MDP) API

Must define:

- state space (S)
- action space (A)
- step (Tr , R , terminated)

Coding

In this course, we will implemented MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Gymnasium provides an **environment** (MDP) API

Must define:

- state space (S)
- action space (A)
- step (Tr, R , terminated)
- reset (s_0)

Coding

In this course, we will implemented MDPs using **gymnasium**

Originally developed by OpenAI for reinforcement learning

Gymnasium provides an **environment** (MDP) API

Must define:

- state space (S)
- action space (A)
- step (Tr , R , terminated)
- reset (s_0)

<https://gymnasium.farama.org/api/env/>

Coding

Gymnasium uses **observations** instead of **states**

Coding

Gymnasium uses **observations** instead of **states**

Question: What was the Markov condition for MDPs?

Coding

Gymnasium uses **observations** instead of **states**

Question: What was the Markov condition for MDPs?

The next Markov state only depends on the current Markov state

Coding

Gymnasium uses **observations** instead of **states**

Question: What was the Markov condition for MDPs?

The next Markov state only depends on the current Markov state

$$\Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_1) = \Pr(s_{t+1} \mid s_t)$$

Coding

Gymnasium uses **observations** instead of **states**

Question: What was the Markov condition for MDPs?

The next Markov state only depends on the current Markov state

$$\Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_1) = \Pr(s_{t+1} \mid s_t)$$

If the Markov property is broken, $s_t \in S$ is not a Markov state

Coding

Gymnasium uses **observations** instead of **states**

Question: What was the Markov condition for MDPs?

The next Markov state only depends on the current Markov state

$$\Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_1) = \Pr(s_{t+1} \mid s_t)$$

If the Markov property is broken, $s_t \in S$ is not a Markov state

Then, we change $s_t \in S$ to an **observation** $o_t \in O$ (more later)

Coding

Gymnasium uses **observations** instead of **states**

Question: What was the Markov condition for MDPs?

The next Markov state only depends on the current Markov state

$$\Pr(s_{t+1} \mid s_t, s_{t-1}, \dots, s_1) = \Pr(s_{t+1} \mid s_t)$$

If the Markov property is broken, $s_t \in S$ is not a Markov state

Then, we change $s_t \in S$ to an **observation** $o_t \in O$ (more later)

Gymnasium uses observations, but for MDPs we treat them as states

Coding

```
import gymnasium as gym

MyMDP(gym.Env):
    def __init__(self):
        self.action_space = gym.spaces.Discrete(3) # A
        self.observation_space = gym.spaces.Discrete(5) # S

    def reset(self, seed=None) -> Tuple[Observation, Dict]

    def step(self, action) -> Tuple[
        Observation, Reward, Terminated, Truncated, Dict
    ]
```

Coding

<https://colab.research.google.com/drive/1rDNik5oRl27si8wdtMLE7Y41U5J2bx-I#scrollTo=9pOLl5OgKvoE>

Exam Next Class

Exam Next Class

We will have an exam next week

Exam Next Class

We will have an exam next week

1 hour 15 minutes, no coding, only math

Exam Next Class

We will have an exam next week

1 hour 15 minutes, no coding, only math

No book, no notes, no calculator – only pencil and pen

Exam Next Class

We will have an exam next week

1 hour 15 minutes, no coding, only math

No book, no notes, no calculator – only pencil and pen

Study notation, probability, bandits, and Markov processes

Exam Next Class

We will have an exam next week

1 hour 15 minutes, no coding, only math

No book, no notes, no calculator – only pencil and pen

Study notation, probability, bandits, and Markov processes

Practice expectations, bandit problems, state transitions, and returns

Exam Next Class

We will have an exam next week

1 hour 15 minutes, no coding, only math

No book, no notes, no calculator – only pencil and pen

Study notation, probability, bandits, and Markov processes

Practice expectations, bandit problems, state transitions, and returns

You must have intuition, not memorize

Exam Next Class

We will have an exam next week

1 hour 15 minutes, no coding, only math

No book, no notes, no calculator – only pencil and pen

Study notation, probability, bandits, and Markov processes

Practice expectations, bandit problems, state transitions, and returns

You must have intuition, not memorize

Too many A's last term, exam will be **difficult**